

# POST-IMPLEMENTATION SECURITY REPORT

**Subject:** Distributed Intrusion Detection System using Snort, Kafka, and ELK

**Date:** 28th December 2025

**Prepared By:** Jyothi Samhitha Katragadda

**Classification:** Demonstration Document

**Intended Audience:** Security Analysts, SOC Engineers, Evaluators

---

## 1. EXECUTIVE SUMMARY

This project implements a **distributed network intrusion detection and monitoring system** using **Snort 3 sensors**, **Apache Kafka**, and the **Elastic Stack (Logstash, Elasticsearch, Kibana)**. The system is designed to monitor **multiple network segments** using independent sensors and forward alerts to a centralized analytics platform in real time.

Two Snort sensors were deployed on separate network interfaces, each monitoring a different segment. Alerts generated by Snort are forwarded through Kafka using custom Python forwarders, ingested by Logstash, indexed into Elasticsearch, and visualized using Kibana dashboards.

The project demonstrates:

- Multi-segment intrusion detection
- Decoupled alert ingestion using Kafka
- Centralized security visibility via ELK
- Practical SOC-style monitoring workflows

## 2. PROJECT METHODOLOGY

### Phase I: Environment Setup

- Windows host system with VirtualBox
- WSL (Ubuntu) for Kafka, Logstash, Elasticsearch, Kibana
- Kali Linux virtual machines for sensors and traffic generation

### Phase II: Snort Sensor Configuration

- Snort 3 installed on Kali Linux
- Custom ICMP detection rules added
- JSON alert output enabled

- Separate log directories per interface

### **Phase III: Kafka & Logstash Pipeline**

- Apache Kafka deployed on WSL & topic created for Snort alerts
- Custom Python forwarder tails Snort alert files
- Logstash pipeline ingests Kafka messages into Elasticsearch

### **Phase IV: Multi-Segment Sensor Deployment**

- Sensor 1: Host-only network (eth0)
- Sensor 2: Internal network (seg-b, eth1)
- Independent Snort processes and forwarders per interface

### **Phase V: Visualization & Validation**

- Elasticsearch index patterns created
- Kibana dashboards built
- Live alert validation through traffic generation
- End-to-end alert flow verified

## **3. THREAT MODEL AND ASSUMPTIONS**

### **Threats Simulated**

- ICMP reconnaissance (ping sweeps)
- TCP SYN scanning
- Network probing activity

### **Assumptions**

- Lab environment is trusted and controlled
- No encrypted payload inspection
- Alerts represent simulated attack behavior
- Sensors operate in passive IDS mode

### **Trust Boundaries**

- Snort sensors generate alerts
- Kafka acts as a trusted transport layer
- ELK stack serves as centralized analysis plane

## 4. SYSTEM ARCHITECTURE

### Core Components

- **Snort 3:** Network intrusion detection
- **Kafka:** Decoupled alert streaming
- **Logstash:** Data ingestion and transformation
- **Elasticsearch:** Alert indexing and storage
- **Kibana:** Visualization and analysis

### Design Rationale

- Kafka prevents alert loss and allows scalability
- Independent sensors allow segment-specific monitoring
- Centralized ELK enables SOC-style analytics

The architecture supports **horizontal sensor scaling** and **real-time alert analysis**.

## 5. QUANTITATIVE ANALYSIS AND METRICS

Metric	Observation
Alerts Generated	ICMP and TCP events detected
Sensors Active	2
Network Segments	Host-only, Internal
Kafka Throughput	Real-time alert delivery
Elasticsearch Docs	Incremental growth verified
Visualization Latency	Near real-time

Document counts increased immediately upon traffic generation, validating the ingestion pipeline.

## 6. OBSERVATIONS AND TECHNICAL ANALYSIS

### Sensor Behavior

- Host-only network generated consistent ICMP alerts
- Internal network traffic required external VM for visibility
- Self-ping traffic does not traverse interfaces

### Pipeline Behavior

- Kafka successfully buffered and forwarded alerts
- Logstash downtime resulted in ingestion gaps
- Elasticsearch indexing confirmed via \_cat/indices

### Visualization Insights

- Clear alert peaks during active scans
- Protocol distribution dominated by ICMP
- Sensor activity validated dual-sensor operation

## 7. CHALLENGES FACED (STAR ANALOGY)

### Situation

Initial alerts were not appearing in Kibana despite traffic generation.

### Task

Identify whether the issue existed at Snort, Kafka, Logstash, or Elasticsearch.

### Action

- Verified Snort logs directly
- Tested Kafka connectivity
- Checked Logstash service status
- Restarted broken ingestion pipeline

### Result

Alert flow restored end-to-end with real-time visualization.

Other challenges included:

- Snort self-traffic misunderstanding
- Python package restrictions in Kali
- Permissions on Snort log files
- Kafka topic persistence confusion

## 8. GAP ANALYSIS: REMAINING RISKS AND MITIGATION

Gap	Risk	Mitigation
No TLS	Data exposure	Enable Kafka TLS
No Auth	Unauthorized access	Add SASL/Auth
Limited Rules	Missed attacks	Expand rule sets
Manual Forwarder	Maintenance	Use Filebeat

## 9. FURTHER IMPROVEMENTS AND ROADMAP

- Add Suricata comparison sensor
- Enable encrypted Kafka transport
- MITRE ATT&CK tagging automation
- Integrate Filebeat instead of Python
- Add anomaly-based detection models
- Long-term alert retention policies

## 10. MITRE ATT&CK MAPPING TABLE

Technique ID	Name	Observed
T1046	Network Service Scanning	Yes
T1018	Network Discovery	Yes
T1040	Network Sniffing	Simulated
T1595	Active Scanning	Yes

## CONCLUSION

This project successfully demonstrates a **distributed intrusion detection architecture** using industry-relevant tools. The system validates real-time detection, scalable ingestion, and centralized security analytics across multiple network segments.

The implementation reflects **SOC-style operational workflows** and provides a strong foundation for further enterprise-grade security monitoring enhancements.